

# Artificial Intelligence

- Artificial Intelligence: A Modern Approach, Stuart J. Russell and Peter Norvig

Artificial Intelligence: A Modern Approach
Part-01: Artificial Intelligence
<b>Introduction</b>
What is AI?
The Foundations of Artificial Intelligence
The History of Artificial Intelligence
The State of the Art
Summary, Bibliographical and Historical Notes, Exercises
<b>Intelligent Agents</b>
Agents and Environments
Good Behavior: The Concept of Rationality
The Name Nature of Environments
The Structure of Agents
Summary, Bibliographical and Historical Notes, Exercises
<b>Problem-Solving</b>
<b>Solving Problems by Searching</b>
Problem-Solving Agents
Example Problems
Searching for Solutions
Uninformed Search Strategies
Informed (Heuristic) Search Strategies
Heuristic Functions
Summary, Bibliographical and Historical Notes, Exercises
<b>Beyond Classical Search</b>
Local Search Algorithms and Optimization Problems
Local Search in Continuous Spaces
Searching with Nondeterministic Actions
Searching with Partial Observations
Online Search Agents and Unknown Environments
Summary, Bibliographical and Historical Notes, Exercises
<b>Adversarial Search</b>
Games
Optimal Decisions in Games
Alpha-Beta Pruning
Imperfect Real-Time Decisions
Stochastic Games
Partially Observable Games
Stage-of-the-Art Game Programs
Alternative Approaches
Summary, Bibliographical and Historical Notes, Exercises
<b>Constraint Satisfaction Problems</b>
Defining Constraint Satisfaction Problems
Constraint Propagation: Inference in CSPs
Backtracking Search for CSPs
Local Search for CSPs
The Structure of Problems

	Summary, Bibliographical and Historical Notes, Exercises
	<b>Part-03: Knowledge, Reasoning, and Planning</b>
	<b>Logical Agents</b>
	Knowledge-Based Agents
	The Wumpus Words
	Logic
	Propositional Logic: A Very Simple Logic
	Propositional Theorem Proving
	Effective Propositional Model Checking
	Agents Based on Propositional Logic
	Summary, Bibliographical and Historical Notes, Exercises
	<b>First-Order Logic</b>
	Representation Revisited
	Syntax and Semantics of First-Order Logic
	Using First-Order Logic
	Knowledge Engineering in First-Order Logic
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Inference in First-Order Logic</b>
	Propositional vs. First-Order Inference
	Unification and Lifting
	Forward Chaining
	Backward Chaining
	Resolutions
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Classical Planning</b>
	Definition of Classical planning
	Algorithms for Planning as State-Space Search
	Planning Graphs
	Other Classical Planning Approaches
	Analysis of Planning Approaches
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Planning and Acting in the Real World</b>
	Time, Schedules, and Resources
	Hierarchical Planning
	Planning and Acting in Nondeterministic Domains
	Multiagent Planning
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Knowledge Representation</b>
	Ontological Engineering
	Categories and Objects
	Events
	Mental Events and Mental Objects
	Reasoning Systems for Categories
	Reasoning with Default Information
	The Internet Shopping World
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Part-04: Uncertain Knowledge and Reasoning</b>
	<b>Quantifying Uncertainty</b>
	Acting Under Uncertainty
	Basic Probability Notation
	Inference Using Full Joint Distributions

	Independence
	Bayes' Rule and Its Use
	The Wumpus World Revisited
	Summary, bibliographical and Historical Notes, Exercises
	<b>Probabilistic Reasoning</b>
	Representing Knowledge in an Uncertain Domain
	The Semantics of Bayesian Networks
	Efficient Representation of Conditional Distributions
	Exact Inference in Bayesian Networks
	Approximate Inference in Bayesian Networks
	Relational and First-Order Probability Models
	Other Approaches to Uncertain Reasoning
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Probabilistic Reasoning Over Time</b>
	Time and Uncertainty
	Inference in temporal Models
	Hidden Markov Models
	Kalman Filters
	Dynamic Bayesian Networks
	Keeping Track of Many Objects
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Making Simple Decisions</b>
	Combining Beliefs and Desires Under Uncertainty
	The Basis of Utility Theory
	Utility Functions
	Multiattribute Utility Functions
	Decision Networks
	The Value of Information
	Decision-Theoretic Expert Systems
	Summary,   Bibliographical and Historical Notes, Exercises
	<b>Making Complex Decisions</b>
	Sequential Decision Problems
	Value Iteration
	Policy Iteration
	Partially Observable MDPs
	Decisions with Multiple Agents: Game Theory
	Mechanism Design
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Part-05: Learning</b>
	<b>Learning from Examples</b>
	Forms of Learning
	Supervised Learning
	Learning Decision Trees
	Evaluating and Choosing the Best Hypothesis
	The Theory of Learning
	Regression and Classification with Linear Models
	Artificial Neural Networks
	Nonparametric Models
	Support Vector Machines
	Ensemble Learning
	Practical Machine Learning

	Summary, Bibliographical and Historical Notes, exercises
	<b>Knowledge in Learning</b>
	A Logical Formulation of Learning
	Knowledge in Learning
	Explanation-Based Learning
	Learning Using Relevance Information
	Inductive Logic Programming
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Learning Probabilistic Models</b>
	Statistical Learning
	Learning with Complete Data
	Learning with Hidden Variables: The EM Algorithm
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Reinforcement Learning</b>
	Introduction
	Passive Reinforcement Learning
	Active Reinforcement Learning
	Generalization in Reinforcement Learning
	Policy Search
	Applications of Reinforcement Learning
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Part-6: Communication, Perceiving, and Acting</b>
	<b>Natural Language Processing</b>
	Language Models
	Text Classification
	Information Retrieval
	Information Extraction
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Natural Language for Communication</b>
	Phrase Structure Grammars
	Syntactic Analysis (Parsing)
	Augmented Grammars and Semantic Interpretation
	Machine Translation
	Speech Recognition
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Perception</b>
	Image Formation
	Early Image-Processing Operations
	Object Recognition by Appearance
	Reconstructing the 3D World
	Object Recognition from Structural Information
	Using Vision
	Summary, bibliographical and Historical Notes, Exercises
	<b>Robotics</b>
	Introduction
	Robot Hardware
	Robotic Perception
	Planning to Move
	Planning Uncertain Movements
	Moving
	Robotic Software Architectures

	Application Domains
	Summary, Bibliographical and Historical Notes, Exercises
	<b>Part-07: Conclusions</b>
	<b>Philosophical Foundations</b>
	Weak AI: Can Machines Act Intelligently?
	Strong AI: Can Machines Really Think?
	The Ethics and Risks of Developing Artificial Intelligence
	Summary, Bibliographical and Historical Notes, Exercises
	<b>AI: The Present and Future</b>
	Agent Components
	Agent Architectures
	Are We Going in the Right Direction?
	What If AI Does Succeed?
	<b>Mathematical Background</b>
	Complexity Analysis and $O()$ Notation
	Vectors, Matrices, and Linear Algebra
	Probability Distributions
	<b>Notes on Languages and Algorithms</b>
	Defining Languages with Backus-Naur Form (BNF)
	Describing Algorithms with Pseudocode
	Online Help

# Machine Learning

- Machine Learning Refined: Foundations, Algorithms, and Applications, Jeremy Watt, Reza Borhani
- Deep Learning, Ian Goodfellow, Yoshua Bengio, Aaron Courville
- Hands-on Machine Learning with Scikit-Learn, Keras & TensorFlow: Concepts, Tools and Techniques to Build Intelligent System, Aurelien Geron

Machine Learning Refined: Foundations, Algorithms, and Applications
<b>Introduction</b>
Teaching a Computer to Distinguish Cats from Dogs – Topics 1
Predictive Learning Problems – Topics 2
Feature Design
Numerical Optimization
Summary
<b>Part-01: Fundamental Tools and Concepts</b>
<b>Fundamentals of Numerical Optimization</b>
Calculus-Defined Optimality – Topics 3
Numerical Methods for Optimization – Topics 4
Summary
Exercises
<b>Regression</b>
The Basics of Linear Regression – Topics 5
Knowledge-Driven Feature Design for Regression – Topics 1
Nonlinear Regression and E2 Regularization – Topics 2
Summary
Exercises
<b>Classification</b>
The Perceptron Cost Functions – Topics 4
The Perceptron Cost Functions – Topics 4
The Logistic Regression Perspective on the Softmax Cost – Topics 2
The Support Vector Machine Perspective on the Margin Perceptron – Topics 4
Multiclass Classification – Topics 4
Knowledge-Driven Feature Design for Classification – Topics 1
Histogram Features for Real Data Types – Topics 3
Summary
Exercises
<b>Part-02: Tools for Fully Data-Dirven Machine Learning</b>
<b>Automatic Feature Design for Regression</b>
Automatic Feature Design for the Ideal Regression Scenario – Topics 3
Automatic Feature Design for the Ideal Regression Scenario – Topics 3
Automatic Feature Design for the Real Regression Scenario – Topics 2
Cross-Validation for Regression – Topics 4
Which Basis Works Best? – Topics 3
Summary
Exercises
Notes on Continuous Function Approximation
<b>Automatic Feature Design for Classification</b>
Automatic Feature Design for the Ideal Classification Scenario – Topics 4
Automatic Feature Design for the Real Classification Scenario – Topics 3
Multiclass Classification – Topics 2

	Cross-Validation for Classification – Topics 4
	Which Basis Works Best?
	Summary
	Exercises
	<b>Kernels, Backpropagation, and Regularized Cross-Validation</b>
	Fixed Feature Kernels – Topics 5
	The Backpropagation Algorithm – Topics 3
	Cross-Validation via L2 Regularization – Topics 3
	Summary
	Further Kernel Calculations – Topics 3
	<b>Part-03: Methods for Large Scale Machine Learning</b>
	<b>Advanced Gradient Schemes</b>
	Fixed Step Length Rules for Gradient Descent – Topics 3
	Adaptive Step Length Rules for Gradient Descent – Topics 2
	Stochastic Gradient Descent – Topics 5
	Convergence Proofs for Gradient Descent Schemes – Topics 4
	Calculation of Computable Lipschitz Constants
	Summary
	Exercises
	<b>Dimension Reduction Techniques</b>
	Techniques for Data Dimension Reduction – Topics 3
	Principal Component Analysis – Topics 1
	Recommender Systems – Topics 2
	Summary
	Exercises
	<b>Part-04: Appendices</b>
	<b>Basic vector and Matrix Operations</b>
	Vector Operations
	Matrix Operations
	<b>Basic of Vector Calculus</b>
	Basic Definitions
	Commonly Used Rules for Computing Derivative
	Examples of Gradient and Hessian Calculations
	<b>Fundamental Matrix Factorizations and The Pseudo-Inverse</b>
	Fundamental Matrix Factorizations – Topics 3
	<b>Convex Geometry</b>
	Definitions of Convexity – Topics 2
	<b>Deep Learning</b>
	<b>Introduction</b>
	Who Should Read This Book?
	Historical trends in Deep Learning
	<b>Part-01: Applied Math and Machine Learning Basics</b>
	<b>Linear Algebra</b>
	Scalars, vectors, Matrices and Tensors
	Multiplying Matrices and Vectors
	Identity and Inverse Matrices
	Linear Dependence and Span
	Norms
	Special Kinds of Matrices and Vectors
	Enigendecomposition
	Singular Value Decomposition

	The Moore-Penrose Pseudoinverse
	The Trace Operator
	The Determinant
	Example: Principal Components Analysis
	<b>Probability and Information Theory</b>
	Why Probability?
	Random Variables
	Probability Distributions
	Marginal Probability
	Conditional Probability
	The Chain Rule of Conditional Probabilities
	Independence and Conditional Independence
	Expectation, Variance and Covariance
	Common Probability Distributions
	Useful Properties of Continuous Variables
	Information Theory
	Structured Probabilistic Models
	<b>Numerical Computation</b>
	Overflow and Underflow
	Poor Conditioning
	Gradient-Based Optimization
	Constrained Optimization
	Example: Linear Least Squares
	<b>Machine Learning Basics</b>
	Learning Algorithms
	Capacity, Overfitting and Under-fitting
	Hyper-parameters and Validation Sets
	Estimators, Bias and Variance
	Maximum Likelihood Estimation
	Bayesian Statistics
	Supervised Learning Algorithms
	Unsupervised Learning Algorithms
	Stochastic Gradient Descent
	Building a Machine Learning Algorithm
	Challenges Motivation Deep Learning
	<b>Part-02: Deep Networks: Modern Practices</b>
	<b>Deep Feedforward Networks</b>
	Example: Learning XOR
	Gradient-Based Learning
	Hidden Units
	Architecture Design
	Back-Propagation and Other Differentiation Algorithms
	Historical Notes
	<b>Regularization for Deep Learning</b>
	Parameter Norm Penalties
	Norm Penalties as Constrained Optimization
	Regularization and Under-Constrained Problems
	Dataset Augmentation
	Noise Robustness
	Semi-Supervised Learning
	Multi-Task Learning



	Early Stopping
	Parameter Tying and Parameter Sharing
	Sparse Representations
	Bagging and Other Ensemble Method
	Dropout
	Adversarial Training
	Tangent Distance, Tangent Prop, and Manifold Tangent Classifier
	<b>Optimization for Training Deep Models</b>
	How Learning Differs from Pure Optimization
	Challenges in Neural Network Optimization
	Basic Algorithms
	Parameter Initialization Strategies
	Algorithms with Adaptive Learning Rates
	Approximate Second-Order Methods
	Optimization Strategies and Meta-Algorithms
	<b>Convolutional Networks</b>
	The Convolution Operation
	Motivation
	Pooling
	Convolution and Pooling as an Infinitely Strong Prior
	Variants for the Basic Convolution Function
	Structured Outputs
	Data Types
	Efficient Convolution Algorithms
	Random or Unsupervised Features
	The Neuroscientific Basis for Convolutional Networks
	Convolutional Network and the History of Deep Learning
	<b>Sequence Modeling: Recurrent and Recursive Nets</b>
	Unfolding Computational Graphs
	Recurrent Neural Networks
	Bidirectional RNNs
	Encoder-Decoder Sequence-to-Sequence Architectures
	Deep Recurrent Networks
	The Challenge of Long-Term Dependencies
	Echo State Networks
	Leaky Units and Other Strategies for Multiple Time Scales
	The Long Short-Term Memory and Other Gated RNNs
	Optimization for Long-Term Dependencies
	Explicit Memory
	<b>Practical Methodology</b>
	Performance Metrics
	Default Baseline Models
	Determining Whether to Gather More Data
	Selecting Hyper-parameters
	Debugging Strategies
	Example: Multi-Digit Number Recognition
	<b>Applications</b>
	Large-Scale Deep Learning
	Computer Vision
	Speech Recognition
	Natural Language Processin

	Other Applications
	<b>Part-03: Deep Learning Research</b>
	<b>Linear Factor Models</b>
	Probabilistic PCA and Factor Analysis
	Independ Component Analysis (ICA)
	Show Feature Analysis
	Sparse Coding
	Manifold Interpretation of PCA
	<b>Auto-encoder</b>
	Under-complete Auto-encoders
	Regularized Auto-encoders
	Representational Power, Layer Size and Depth
	Stochastic Encoders and Decoders
	Denoising Auto-encoders
	Learning Manifolds with Auto-encoders
	Contractive Auto-encoders
	Predictive Sparse Decomposition
	Applications of Auto-encoders
	<b>Representation Learning</b>
	Greedy Layer-Wise Unsupervised Pre-training
	Transfer Learning and Domain Adaptation
	Semi-Supervised Disentangling of Causal Factors
	Distributed Representation
	Exponential Gains from Depth
	Providing Clues to Discover Underlying Causes
	<b>Structured Probabilistic Models for Deep Learning</b>
	The Challenge of Unstructured Modeling
	Using Graphs to Describe Model Structure
	Sampling from Graphical Models
	Advantages of Structured Modeling
	Learning about Dependencies
	Inference and Approximate Inference
	The Deep Learning Approach to Structured Probabilistic Models
	<b>Monte Carlo Methods</b>
	Sampling and Monte Carlo Methods
	Importance Sampling
	Markov Chain Monte Carlo Methods
	Gibbs Sampling
	The Challenge of Mixing Between Separated Modes
	<b>Confronting the Partition Function</b>
	The Log-Likelihood Gradient
	Stochastic Maximum Likelihood and Contrastive Divergence
	Pseudolikelihood
	Score Matching and Ratio Matching
	Denoising Score Matching
	Noise-Contrastive Estimation
	Estimating the Partition Function
	<b>Approximate Inference</b>
	Inference as Optimization
	Expectation Maximization
	MAP Inference and Sparse Coding

	Variational Inference and Learning
	Learned Approximate Inference
	<b>Deep Generative Models</b>
	Boltzmann Machines
	Restricted Boltzmann Machines
	Deep Belief Networks
	Deep Boltzmann Machines
	Boltzmann Machines for Real-Valued Data
	Convolutional Boltzmann Machines
	Boltzmann Machines for Structured or Sequential Outputs
	Other Boltzmann Machines
	Back-Propagation Through Random Operations
	Directed Generative Nets
	Drawing Samples from Auto-encoders
	Generative Stochastic Networks
	Other Generation Schemes
	Evaluating Generative Models
	Conclusion
	<b>Hands-on Machine Learning with Scikit-Learn, Keras &amp; TensorFlow</b>
	<b>Part-01: The Fundamentals of Machine Learning</b>
	<b>The Machine Learning Landscape</b>
	What is Machine Learning?
	Why Use Machine Learning?
	Types of Machine Learning Systems – Topics 3
	Main Challenges of Machine Learning – Topics 7
	Testing and Validating – Topics 2
	Exercise
	<b>End-to-End Machine Learning Project</b>
	Working with Real Data
	Look at the Big Picture – Topics 3
	Get the Data – Topics 4
	Discover and Visualize the Data to Gain Insights – Topics 3
	Prepare the Data for Machine Learning Algorithms – Topics 5
	Select and Train a Model – Topics 2
	Fine-Tune Your Model – Topics 5
	Launch, Monitor, and Maintain Your System
	Try It Out!
	Exercises
	<b>Classification</b>
	MNIST
	Training a Binary Classifier
	Performance Measures – Topics 5
	Multiclass Classification
	Error Analysis
	Multilevel Classification
	Multi-output Classification
	Exercises
	<b>Training Models</b>
	Linear Regression – Topics 2
	Gradient Descent – Topics 3
	Polynomial regression

	Learning Curves
	Regularized Linear Models – Topics 4
	Logistic Regression – Topics 4
	Exercises
	<b>Support Vector Machines</b>
	Linear SVM Classification   Soft Margin Classification
	Nonlinear SVM Classification – Topics 4
	SVM Regression
	Under the Hood – Topics 6
	Exercise
	<b>Decision Trees</b>
	Training and Visualizing a Decision Tree
	Making Predictions
	Estimating Class Probabilities
	The CART Training Algorithm
	Computational Complexity
	Gini Impurity or Entropy?
	Regularization Hyper-parameters
	Regression
	Instability
	Exercises
	<b>Ensemble Learning and Random Forests</b>
	Voting Classifiers
	Bagging and Pasting – Topics 2
	Random Patches and Random Subspaces
	Random Forests – Topics 2
	Boosting – Topics 2
	Stacking
	Exercises
	<b>Dimensionality Reduction</b>
	The Curse of Dimensionality
	Main Approaches for Dimensionality Reduction – Topics 2
	PCA – Topics 9
	Kernel PCA   Selecting a Kernel and Tuning Hyper-parameters
	LLE
	Other Dimensionality Reduction Techniques
	Exercises
	<b>Unsupervised Learning Techniques</b>
	Clustering – Topics 7
	Gaussian Mixtures – Topics 4
	<b>Part-02: Neural Networks and Deep Learning</b>
	<b>Introduction to Artificial Neural Networks with Keras</b>
	From Biological to Artificial Neurons – Topics 6
	Implementing MLPs with Keras – Topics 8
	Fine-Tuning Neural Network hyper-parameters – Topics 3
	Exercises
	<b>Training Deep Neural Networks</b>
	Vanishing/Exploding Gradients Problems – Topics 4
	Reusing Pretrained Layers – Topics 3
	Faster Optimizers – Topics 6
	Avoiding Overfitting Through Regularization – Topics 4

	Summary and Practical Guidelines
	Exercises
	<b>Custom Models and Training with TensorFlow</b>
	A Quick Tour of Tensor Flow
	Using TensorFlow like NumPy – Topics 5
	Customizing Models and Training Algorithms – Topics 9
	TensorFlow Functions and Graphs – Topics 2
	<b>Loading and Preprocessing Data with TensorFlow</b>
	The Data API – Topics 6
	The TFRecord Format – Topics 5
	The Features API – Topics 6
	TF Transform
	The TensorFlow Datasets (TFDS) Project
	<b>Deep Computer Vision Using Convolutional Neural Networks</b>
	The Architecture of the Visual Cortex
	Convolutional Layer – Topics 4
	Pooling Layer   TensorFlow Implementation
	CNN Architectures – Topics 7
	Implementing a ResNet-34 CN Using Keras
	Using Pretrained Models From Keras
	Pretrained Models for Transfer Learning
	Classification and Localization
	Object Detection – Topics 2
	Semantic Segmentation
	Exercises

# Data Science

- Introducing Data Science: Big Data, Machine Learning, and More, Using Python Tools, Davy Cielen, Arno D.B.
- Data Science From Scratch: First Principles with Python, Joel Grus

Introducing Data Science: Big Data, Machine Learning and More
<b>Data Science in a big Data World</b>
Benefits and Uses of Data Science and Big Data
Facets of Data – Topics 7
The Data Science Process – Topics 6
The Big Data Ecosystem and Data Science – Topics 11
An Introductory Working Example of Hadoop
Summary
<b>The Data Science Process</b>
Overview of the Data Science Process – Topics 1
Step-1: Defining Research Goals and Creating a Project Charter – Topics 2
Step-2: Retrieving Data – Topics 3
Step-3: Cleansing, integrating and Transforming Data – Topics 4
Step-4: Exploratory Data Analysis
Step-5: Build the Models – Topics 3
Step-6: Presenting Findings and Building Applications on top of Them
Summary
<b>Machine Learning</b>
What is Machine Learning and Why Should You Care About It? – Topics 3
The Modeling Process – Topics 4
Types of Machine Learning – Topics 2
Semi-supervised Learning
Summary
<b>Handling Large Data on a Single Computer</b>
The Problems You Face When Handling Large Data
General Techniques for Handling Large Volumes of Data – Topics 3
General Programming Tips for Dealing with Large Data Sets – Topics 4
Case Study-1: Predicting Malicious URLs – Topics 4
Case Study-2: Building a Recommender System Inside a Database – Topics 5
Summary
<b>First Steps in Big Data</b>
Distributing Data Storage and Processing with Frameworks – Topics 2
Case Study: Assessing Risk When Loaning Money – Topics 5
Summary
<b>Join the NoSQL Movement</b>
Introduction to NoSQL – 4
Case Study: What Disease is That? – 6
Summary
<b>The Rise of Graph Databases</b>
Introducing Connected Data and Graph Databases – Topics 1
Introducing Neo4J: A Graph Database – Topics 1
Connected Data Example: A Recipe Recommendation Engine – Topics 6
Summary
<b>Text Mining and Text Analytics</b>
Text Mining in the Real World

	Text Mining Techniques – Topics 3
	Case Study: Classifying Reddit Posts – Topics 9
	Summary
	<b>Data Visualization to the End User</b>
	Data Visualization Options
	Cross-filter, The JavaScript MapReduce Library – Topics 2
	Creating an Interactive Dashboard with Dc.js
	Dashboard Development Tools
	Summary
	<b>Data Science From Scratch: First Principles with Python</b>
	<b>Introduction</b>
	The Ascendance of Data
	What is Data Science?
	Motivating Hypothetical: DataSciencester – Topics 6
	<b>A Crash Course in Python</b>
	The Basics – Topics 14
	The Not-So-Basics – Topics 11
	For Further Exploration
	<b>Visualizing Data</b>
	matplotlib
	Bar Charts
	Line Charts
	Scatterplots
	For Further Exploration
	<b>Linear Algebra</b>
	Vectors
	Matrices
	For Further Exploration
	<b>Statistics</b>
	Describing a Single Set of Data – Topics 2
	Correlation
	Simpson's Paradox
	Some Other Correlational Caveats
	Correlation and Causation
	For Further Exploration
	<b>Probability</b>
	Dependency and Independency
	Conditional Probability
	Bayes's Theorem
	Random Variables
	Continuous Distributions
	The Normal Distribution
	The Central Limit Theorem
	For Further Exploration
	<b>Hypothesis and Inference</b>
	Statistical Hypothesis Testing
	Example: Flipping a Coin
	Confidence Intervals
	P-hacking
	Example: Running an A/B Test
	Bayesian Inference

	For Further Exploration
	<b>Gradient Descent</b>
	The Idea Behind Gradient Descent
	Estimating the Gradient
	Using the Gradient
	Choosing the Right Step Size
	Putting It All Together
	Stochastic Gradient Descent
	For Further Exploration
	<b>Getting Data</b>
	Stdin and stdout
	Reading Files – Topics 2
	Scraping the Web – Topics 2
	Using APIs – Topics 3
	Example: Using the Twitter APIs   Getting Credentials
	For Further Exploration
	<b>Working with Data</b>
	Exploring Your Data – Topics 3
	Cleaning and Munging
	Manipulating Data
	Rescaling
	Dimensionality Reduction
	For Further Exploration
	<b>Machine Learning</b>
	Modeling
	What is Machine Learning?
	Overfitting and Under-fitting
	Correctness
	The Bias-Variance Trade-off
	Feature Extraction and Selection
	For Further Exploration
	<b>K-Nearest Neighbors</b>
	The Model
	Example: Favorite Languages
	The Curse of Dimensionality
	For Further Exploration
	<b>Naïve Bayes</b>
	A Really Dumb Spam Filter
	A More Sophisticated Spam Filter
	Implementation
	Testing Our Model
	For Further Exploration
	<b>Simple Linear Regression</b>
	The Model
	Using Gradient Descent
	Maximum Likelihood Estimation
	For Further Exploration
	<b>Multiple Regression</b>
	The Model
	Further Assumptions of the Least Squares Model
	Fitting the Model



	Interpreting the Model
	Goodness of Fit
	Digression: The Bootstrap
	Standard Errors of Regression Coefficients
	Regularization
	For Further Exploration
	<b>Logistic Regression</b>
	The Problem
	The Logistic Function
	Applying the Model
	Goodness of Fit
	Support Vector Machines
	For Further Investigation
	<b>Decision Trees</b>
	What is a Decision Tree?
	Entropy
	The Entropy of a Partition
	Creating a Decision Tree
	Putting It All Together
	Random Forests
	For Further Exploration
	<b>Neural Networks</b>
	Perceptrons
	Feed-Forward Neural Networks
	Backpropagation
	Example: Defeating a CAPTCHA
	For Further Exploration
	<b>Clustering</b>
	The Idea
	The Model
	Example: Meetups
	Choosing K
	Example: Clustering Colors
	Bottom-up Hierarchical Clustering
	For Further Exploration
	<b>Natural Language Processing</b>
	World Clouds
	n-gram Models
	Grammars
	An Aside: Gibbs Sampling
	Topic Modeling
	For Further Exploration
	<b>Network Analysis</b>
	Betweenness Centrality
	Eigenvector Centrality – Topics 2
	Directed Graphs and PageRank
	For Further Exploration
	<b>Recommender Systems</b>
	Manual Curation
	Recommending What's Popular
	User-Based Collaborative Filtering

	Item-Based Collaborative Filtering
	For Further Exploration
	<b>Databases and SQL</b>
	CREATE TABLE and INSERT
	UPDATE
	DELETE
	SELECT
	GROUP BY
	ORDER BY
	JOIN
	Subqueries
	Indexes
	Query Optimization
	NoSQL
	For Further Exploration
	<b>MapReduce</b>
	Example: Word Count
	Why MapReduce
	MapReduce More Generally
	Example: Analyzing Status Updates
	Example: Matrix Multiplication
	An Aside: Combiners
	For Further Exploration
	<b>Go Forth and Do Data Science</b>
	IPython
	Mathematics
	Not from Scratch – Topics 5
	Find Data
	Do Data Science – Topics 4

# Data Warehouse and Mining

- Data Mining Concepts and Techniques, Jiawei Han, Micheline Kamber, Jian Pei
- Building the Data Warehouse, William H. Inmon

<b>Data Mining Concepts and Techniques</b>
<b>Introduction</b>
Why Data Mining? – Topics 2
What is Data Mining?
What Kinds of Data Can Be Mined? – Topics 4
What kinds of Patterns Can Be Mind? – Topics 6
Which Technologies Are Used? – Topics 4
Which Kinds of Applications Are Targeted? – Topics 2
Major Issues in Data Mining – Topics 5
Summary
Exercises
Bibliographic Nots
<b>Getting to Know Your Data</b>
Data Objects and Attribute Types – Topics 6
Basic Statistical Descriptions of Data – Topics 3
Data Visualization – Topics 5
Measuring Data Similarity and Dissimilarity – Topics 7
Summary
Exercises
Bibliographic Notes
<b>Data Processing</b>
Data Preprocessing: An Overview – Topics 2
Data Cleaning – Topics 3
Data Integration – Topics 4
Data Reduction – Topics 9
Data Transformation and Data Discretization – Topics 6
Summary
Exercises
Bibliographic Notes
<b>Data Warehousing and Online Analytical Processing</b>
Data Warehouse: Basic Concepts – Topics 7
Data Warehouse Modeling: Data Cube and OLAP – Topics 6
Data Warehouse Design and Usage – Topics 4
Data Warehouse Implementation – Topics 4
Data Generalization by Attribute-Oriented Induction – Topics 3
Summary
Exercises
Bibliographic Notes
<b>Data Cube Technology</b>
Data Cube Computation: Preliminary Concepts – Topics 2
Data Cube Computation Methods – Topics 4
Processing Advanced Kinds of Queries by Exploring Cube Technology – Topics 2
Multidimensional Data Analysis in Cube Space – Topics 3
Summary
Exercises

	Bibliographic Notes
	<b>Mining Frequent Patterns, Associations, and Correlations: Basic Concepts and Methods</b>
	Basic Concepts – Topics 2
	Frequent Itemset Mining Methods – Topics 6
	Which Patterns Are Interesting? – Pattern Evaluation Methods – Topics 3
	Summary
	Exercises
	Bibliographic Notes
	<b>Advanced Pattern Mining</b>
	Pattern Mining: A Road Map
	Pattern Mining in Multilevel, Multidimensional Space – Topics 4
	Constraint-Based Frequent Pattern Mining – Topics 2
	Mining High-Dimensional Data and Colossal Patterns – Topics 1
	Mining Compressed or Approximate Patterns – Topics 2
	Pattern Exploration and Application – Topics 2
	Summary
	Exercises
	Bibliographic Notes
	<b>Classification: Basic Concepts</b>
	Basic Concepts – Topics 2
	Decision Tree Induction – Topics 5
	Bayes Classification Methods – Topics 2
	Rule-Based Classification – Topics 3
	Model Evaluation and Selection – Topics 6
	Techniques to Improve Classification Accuracy – Topics 5
	Summary
	Exercises
	Bibliographic Notes
	<b>Classification: Advanced Methods</b>
	Bayesian Belief Networks – Topics 2
	Classification by Backpropagation – Topics 4
	Support Vector Machines – Topics 2
	Classification Using Frequent Patterns – Topics 3
	Lazy Learners (or Learning from Your Neighbors) – Topics 2
	Other Classification Methods – Topics 3
	Additional Topics Regarding Classification – Topics 4
	Summary
	Exercises
	Bibliographic Notes
	<b>Cluster Analysis: Basic Concepts and Methods</b>
	Cluster Analysis – Topics 3
	Partitioning Methods – Topics 2
	Hierarchical Method – Topics 5
	Density-Based Methods – Topics 3
	Grid-Based Methods – Topics 2
	Evaluation of Clustering – Topics 3
	Summary Exercises
	Bibliographic Notes
	<b>Advanced Cluster Analysis</b>
	Probabilistic Model-Based Clustering – Topics 3
	Clustering High-Dimensional Data – Topics 4

	Clustering Graph and Network Data – Topics 3
	Clustering with Constraints – Topics 2
	Summary
	Exercises
	Bibliographic Notes
	<b>Outlier Detection</b>
	Outliers and Outlier Analysis – Topics 3
	Outlier Detection Methods – Topics 2
	Statistical Approaches – Topics 2
	Proximity-Based Approaches – Topics 3
	Clustering-Based Approaches
	Classification-Based Approaches
	Mining Contextual and Collective Outliers – Topics 3
	Outlier Detection in High-Dimensional Data – Topics 3
	Summary
	Exercises
	Bibliographic Notes
	<b>Data Mining Trends and Research Frontiers</b>
	Mining Complex Data Types – Topics 3
	Other Methodologies of Data Mining – Topics 3
	Data mining Applications – Topics 5
	Data Mining and Society – Topics 2
	Data Mining Trends
	Summary
	Exercises
	Bibliographic Notes
	<b>Building the Data Warehouse</b>
	<b>Evolution of Decision Support Systems</b>
	The Evolution – Topics 5
	Problems with the Naturally Evolving Architecture – Topics 7
	The Development Life Cycle
	Patterns of Hardware Utilization
	Setting the Stage for Re-Engineering
	Monitoring the Data Warehouse Environment
	Summary
	<b>The Data Warehouse Environment</b>
	The Structure of the Data Warehouse
	Subject Orientation
	Day 1 to Day n Phenomenon
	Granularity – Topics 3
	Exploration and Data Mining
	Living Sample Database
	Partitioning as a Design Approach   Partitioning of Data
	Structuring Data in the Data Warehouse
	Auditing and the Data Warehouse
	Data Homogeneity and Heterogeneity
	Purging Warehouse Data
	Reporting and the Architected Environment
	The Operational Window of Opportunity
	Incorrect Data in the Data Warehouse
	Summary

	<b>The Data Warehouse and Design</b>
	Beginning with Operational Data
	Process and Data Models and the Architected Environment
	The Data Warehouse and Data Models – Topics 3
	The Data Model and Iterative Development
	Normalization and De-normalization – Topics 1
	Metadata – Topics 1
	Cyclicity of Data – The Wrinkle of Time
	Complexity of Transformation and Integration
	Triggering the Data Warehouse Record – Topics 3
	Profile Records
	Managing Volume
	Creating Multiple Profile Records
	Going from the Data Warehouse to the Operational Environment
	Direct Operational Access of Data Warehouse Data
	Indirect Access of Data Warehouse Data – Topics 3
	Indirect Use of Data Warehouse Data
	Start Joins
	Supporting the ODS
	Requirements and the Zachman Framework
	Summary
	<b>Granularity in the Data Warehouse</b>
	Managing Large Amounts of Data
	Managing Multiple Media
	Indexing and Monitoring Data
	Interfaces of Many Technologies
	Programmer or Designer Control of Data Placement
	Parallel Storage and Management of Data   Metadata Management
	Language Interface
	Efficient Loading of Data
	Efficient Index Utilization
	Compaction of Data
	Compound Keys
	Variable-Length Data
	Lock Management
	Index-Only Processing
	Fast Restore
	Other Technological Features
	DBMS Types and the Data Warehouse
	Changing DBMS Technology
	Multidimensional DBMS and the Data Warehouse
	Data Warehousing Across Multiple Storage Media
	The Role of Metadata in the Data Warehouse Environment
	Context and Content – Topics 3
	Refreshing the Data Warehouse
	Testing
	Summary
	<b>The Distributed Data Warehouse</b>
	Types of Distributed Data Warehouses – Topics 8
	The Nature of the Development Efforts – Topics 1
	Distributed Data Warehouse Development – Topics 3

	Building the Warehouse on Multiple Levels
	Multiple Groups Building the Current Level of Detail – Topics 3
	Multiple Platforms for Common Detail Data
	Summary
	<b>Executive Information Systems and the Data Warehouse</b>
	EIS – The Promise
	A Simple Example
	Drill-Down Analysis
	Supporting the Drill-Down Process
	The Data Warehouse as a Basis for EIS
	Where to Turn
	Event Mapping
	Detailed Data and EIS
	Keeping Only Summary Data in the EIS
	Summary
	<b>External Data and the Data Warehouse</b>
	External Data in the Data Warehouse
	Metadata and External Data
	Storing External Data
	Different Components of External Data
	Modeling and External Data
	Secondary Reports
	Archiving External Data
	Comparing Internal Data to External Data
	Summary
	<b>Migration to the Architected Environment</b>
	A Migration Plan
	The Feedback Loop
	Strategic Considerations
	Methodology and Migration
	A Data-Driven Development Methodology – Topics 3
	Summary
	<b>The Data Warehouse and the Web</b>
	Supporting the e-Business Environment
	Moving Data from the Web to the Data Warehouse
	Moving Data from the Data Warehouse of the Web
	Web Support
	Summary
	<b>Unstructured Data and the Data Warehouse</b>
	Integrating the Two Worlds – Topics 5
	A Themed Match – Topics 4
	A Two-Tiered Data Warehouse – Topics 3
	A Self-Organizing Map (SOM) – Topics 2
	Fitting the Two Environments Together
	Summary
	<b>The Really Large Data Warehouse</b>
	Why The Rapid Growth?
	The Impact of Large Volumes of Data – Topics 7
	Disk Storage in the Face of Data Separation – Topics 4
	Moving Data from One Environment to Another – Topics 4
	Inverting the Data Warehouse

	Total Cost
	Maximum Capacity
	Summary
	<b>The Relational and the Multidimensional Models as a Basis for Database Design</b>
	The Relational Model
	The Multidimensional Model
	Snowflake Structures
	Differences Between the Models – Topics 5
	Independent Data Marts
	Building Independent Data Marts
	Summary
	<b>Data Warehouse Advanced Topics</b>
	End-User Requirements and the Data Warehouse – Topics 3
	Resource Contention in the Data Warehouse – Topics 4
	Data Marts and Data Warehouses in the Same Processor
	The Life Cycle of Data – Topics 1
	Testing and the Data Warehouse
	Tracing the Flow of Data Through the Data Warehouse – Topics 2
	Data Warehouse and the Web-Based e-Business Environment – Topics 4
	The Financial Data Warehouse
	The System of Record
	A Brief History of Architecture – Evolving to the Corporate Information Factory – Topics 2
	CIF – Into the Future – Topics 4
	Summary
	<b>Cost-Justification and Return on Investment for a Data Warehouse</b>
	Copying the Competition
	The Macro Level of Cost-Justification
	A Micro Level Cost-Justification
	Information from the Legacy Environment – Topics 6
	The Time Value of Data   The Speed of Information
	Integrated Information – Topics 2
	Summary
	<b>The Data Warehouse and the ODS</b>
	Complementary Structures – Topics 3
	Different Classes of ODS
	Database Design – A Hybrid Approach
	Drawn to Proportion
	Transaction Integrity in the ODS
	Time Slicing the ODS Day
	Multiple ODS
	ODS and the Web Environment
	An Example of an ODS
	Summary
	<b>Corporate Information Compliance and Data Warehousing</b>
	The Basic Activities
	Financial Compliance – Topics 2
	Auditing Corporate Communications
	Summary
	<b>The End-User Community</b>
	The Farmer
	The Explorer



	The Miner
	The Tourist
	The Community
	Different Types of Data
	Cost-Justification and ROI Analysis
	Summary
	<b>Data Warehouse Design Review Checklist</b>
	When to Do a Design Review
	Who Should Be in the Design Review?
	What Should the Agenda Be?
	The Results
	Administering the Review
	A Typical Data Warehouse Design Review
	Summary

# Combinatorial Optimization

➤ Combinatorial Optimization: Theory and Algorithms, Bernhard Korte, Jens Vygen

Combinatorial Optimization
<b>Introduction</b>
Enumeration
Running Time of Algorithms
Linear Optimization Problems
Sorting
Exercises
References
<b>Graphs</b>
Basic Definitions
Trees, Circuits, and Cuts
Connectivity
Eulerian and Bipartite Graphs
Planarity
Planar Duality
Exercises
References
<b>Linear Programming</b>
Polyhedra
The Simplex Algorithm
Duality
Convex Hulls and Polytopes
Exercises
References
<b>Linear Programming Algorithms</b>
Size of Vertices and Faces
Continued Fractions
Gaussian Elimination
The Ellipsoid Method
Khachiyan's Theorem
Separation and Optimization
Exercises
<b>Integer Programming</b>
The Integer Hull of a Polyhedron
Unimodular Transformations
Total Dual Integrality
Totally Unimodular Matrices
Cutting Planes
Lagrangian Relaxation
Exercises
References
<b>Spanning Trees and Arborescences</b>
Minimum Spanning Trees
Minimum Weight Arborescences
Polyhedral Descriptions
Packing Spanning Trees and Arborescences

	Exercises
	References
	<b>Shortest Paths</b>
	Shortest Paths Form One Source
	Shortest Paths Between All Pairs of Vertices
	Minimum Mean Cycles
	Exercises
	References
	<b>Network Flows</b>
	Max-Flow-Min-Cut Theorem
	Mengers Theorem
	The Edmonds-Karp algorithm
	Blocking Flows
	The Goldberg-Tarjan Algorithm
	Gomory-Hu Trees
	The Minimum Cut in an Undirected Graph
	Exercises
	References
	<b>Minimum Cost Flows</b>
	Problem Formulation
	An Optimality Criterion
	Minimum Mean Cycle-Cancelling Algorithm
	Successive Shortest Path Algorithm
	Orlin's Algorithm
	Exercises
	References
	<b>Maximum Matchings</b>
	Bipartite Matching
	The Tuttel Matrix
	Tutte's Theorem
	Ear-Decompositions of Factor-Critical Graphs
	Edmonds' Matching Algorithm
	Exercises
	References
	<b>Weighted Matching</b>
	The Assignment Problem
	Outline of the Weighted Matching Algorithm
	Implementation of the Weighted Matching Algorithm
	Postoptimality
	The Matching Polytope
	Exercises
	References
	<b>b-Matchings and T-Joins</b>
	b-Matchings
	Minimum Weight T-Joins
	T-Joins and T-Cuts
	The Padberg-Rao Theorem
	Exercises
	References
	<b>Matroids</b>
	Independence Systems and Matroids

	Other Matroid Axioms
	Duality
	The Greedy Algorithm
	Matroid Intersection
	Matroid Partitioning
	Weighted Matroid Intersection
	Exercises
	References
	<b>Generalizations of Matroids</b>
	Greedoids
	Polymatroids
	Minimizing Submodular Functions
	Exercises
	References
	<b>NP-Completeness</b>
	Turing Machines
	Church's Thesis
	P and NP
	Cook's Theorem
	Some Basic NP-Complete Problems
	The Class coNp
	NP-Hard Problems
	Exercises
	References
	<b>Approximation Algorithms</b>
	Set Converting
	Colouring
	Approximation Schemes
	Maximum Satisfiability
	The PCP Theorem
	L-Reductions
	Exercises
	References
	<b>The Knapsack Problem</b>
	Fractional Knapsack and Weighted Median Problem
	A Pseudopolynomial Algorithm
	A Full Polynomial Approximation Scheme
	Exercises
	References
	<b>Bin-Packing</b>
	Greedy Heuristics
	An Asymptotic Approximation Scheme
	The Karmarkar-Karp Algorithm
	Exercises
	References
	<b>Multicommodity Flows and Edge-Disjoint Paths</b>
	Multicommodity Flows
	Algorithms for Multicommodity Flows
	Directed Edge-Disjoint Paths Problem
	Undirected Edge-Disjoint Paths Problem
	Exercises

	References
	<b>Network Design Problems</b>
	Steiner Trees
	Survivable Network Design
	A Primal-Dual Approximation Algorithm
	Jain's Algorithm
	Exercises
	References
	<b>The Traveling Salesman Problem</b>
	Approximation Algorithms for the TSP
	Euclidean TSPs
	Local Search
	The Traveling Salesman Polytope
	Lower Bounds
	Branch-and-Bound
	Exercises
	References

# Numerical Analysis

- An Introduction to Numerical Methods and Analysis, James F. Epperson

An Introduction to Numerical Methods and Analysis
<b>Introductory Concepts and Calculus Review</b>
Basic Tools of Calculus – Topics 2
Error, Approximate Equality, and Asymptotic Order Notation – Topics 3
A Primer on Computer Arithmetic
A Word on Computer Language and Software
Simple Approximations
Application: Approximating the Natural Logarithm
A Brief History of Computing
Literature Review
References
<b>A Survey of Simple Methods and Tools</b>
Horner's Rule and Nested Multiplication
Difference Approximations to the Derivative
Application: Euler's Method for Initial Value Problems
Linear Interpolation
Application – The Trapezoid Rule
Solution of Tridiagonal Linear Systems
Application: Simple Two-Point Boundary Value Problems
<b>Root-Finding</b>
The Bisection Method
Newton's Method: Derivation and Examples
How to Stop Newton's Method
Application: Division Using Newton's Method
The Newton Error Formula
Newton's Method: Theory and Convergence
Application: Computation of the Square Root
The Secant Method: Derivation and Examples
Fixed-Point Iteration
Roots of Polynomials, Part 1
Special Topics in Root-finding Methods – Topics 5
Very High-order Methods and the Efficiency Index
Literature and Software Discussion
References
<b>Interpolation and Approximation</b>
Lagrange Interpolation
Newton Interpolation and Divided Differences
Interpolation Error
Application: Muller's Method and Inverse Quadratic Interpolation
Application: More Approximations to the Derivative
Hermite Interpolation
Piecewise Polynomial Interpolation
An Introduction to Splines – Topics 2
Application: Solution of Boundary Value Problems
Tension Splines
Least Squares Concepts in Approximation

	An Introduction to Data Fitting
	Least Squares Approximation and Orthogonal Polynomials
	Advanced Topics in Interpolation Error – Topics 3
	Literature and Software Discussion
	References
	<b>Numerical Integration</b>
	A Review of the Definite Integral
	Improving the Trapezoid Rule
	Simpson's Rule and Degree of Precision
	The Midpoint Rule
	Application: Stirling's Formula
	Gaussian Quadrature
	Extrapolation Methods
	Special Topics in Numerical Integration – Topics 4
	Literature and Software Discussion
	References
	<b>Numerical Methods for Ordinary Differential Equations</b>
	The Initial Value Problem: Background
	Euler's Method
	Analysis of Euler's Method
	Variants of Euler's Method – Topics 4
	Single-Step Methods: Runge-Kutta
	Multistep Methods – Topics 2
	Stability Issues – Topics 2
	Application of Systems of Equations – Topics 2
	Adaptive Solvers
	Boundary Value Problems – Topics 3
	Literature and Software Discussion
	References
	<b>Numerical Methods for the Solution of Systems of Equations</b>
	Linear Algebra Review
	Linear System and Gaussian Elimination
	Operation Counts
	The LU Factorization
	Perturbation, Conditioning, and Stability – Topics 4
	SPD Matrices and the Cholesky Decomposition
	Iterative Methods for Linear Systems: A Brief Survey
	Nonlinear Systems: Newton's Method and Related Ideas – Topics 2
	Application: Numerical Solution of Nonlinear Boundary Value Problems
	Literature and Software Discussion
	References
	<b>Approximate Solution of the Algebraic Eigenvalue Problem</b>
	Eigenvalue Review
	Reduction to Hessenberg Form
	Power Methods
	An Overview of the QR Iteration
	Application: Roots of Polynomials, Part 2
	Literature and Software Discussion
	References
	<b>A Survey of Numerical Methods for Partial Differential Equations</b>
	Difference Methods for the Diffusion Equation – Topics 3

	Finite Element Methods for the Diffusion Equation
	Difference Methods for Poisson Equations – Topics 3
	Literature and Software Discussion
	References
	<b>An Introduction to Spectral Methods</b>
	Spectral Methods for Two-Point Boundary Value Problems
	Spectral Methods for Time-Dependent Problems
	Clenshaw-Curtis Quadrature
	Literature and Software Discussion
	References
	<b>Appendix: Proofs of Selected Theorems, And Additional Material</b>
	Proofs of the Interpolation Error Theorems
	Proof of the Stability Result for ODEs
	Stiff Systems of Differential Equations and Eigenvalues
	The Matrix Perturbation Theorem